



# Running jupyter notebook remotely in a docker swarm cluster

Jordi Deu-Pons  
*Barcelona 2017*

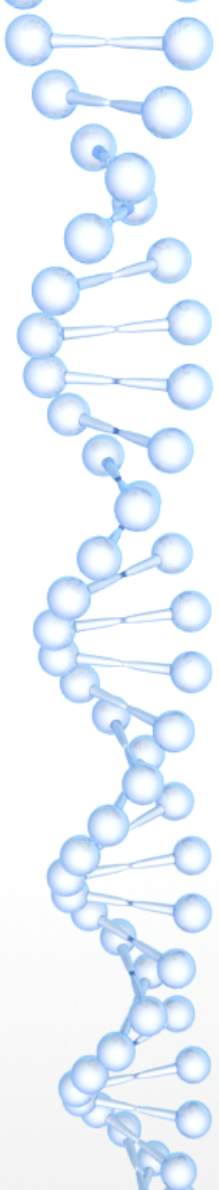


jupyter



docker



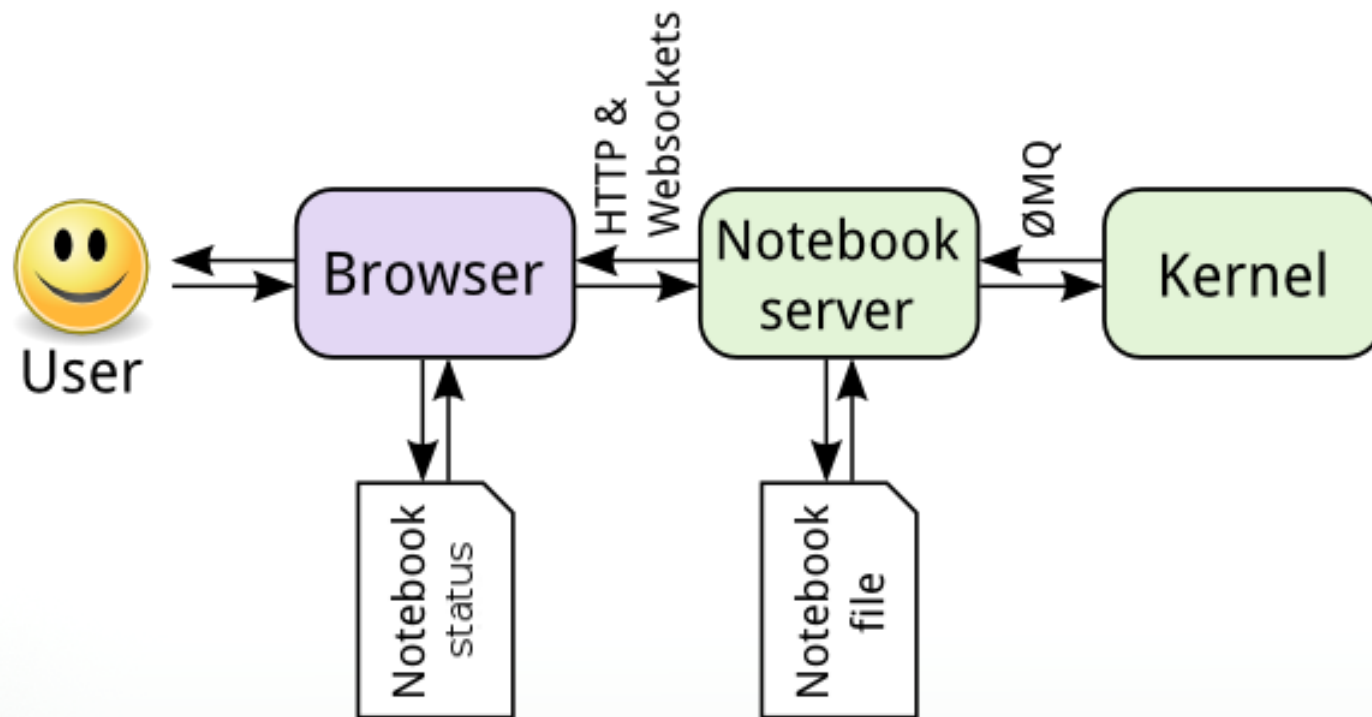


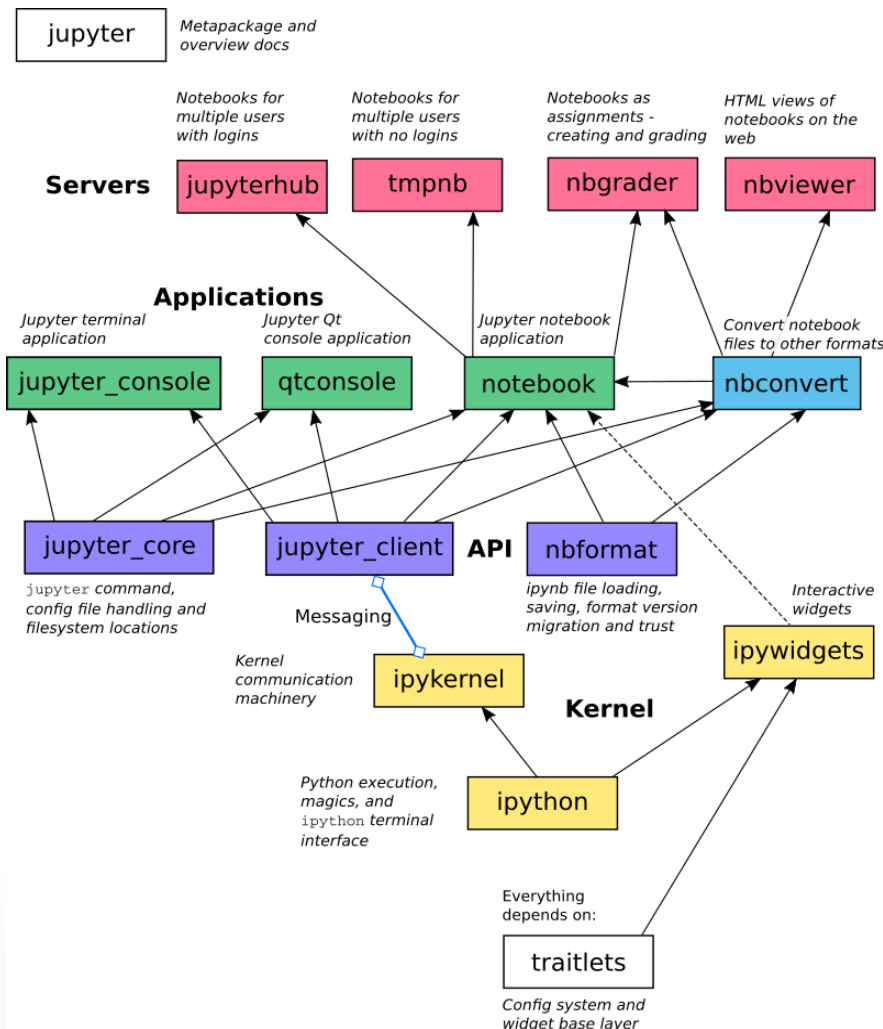
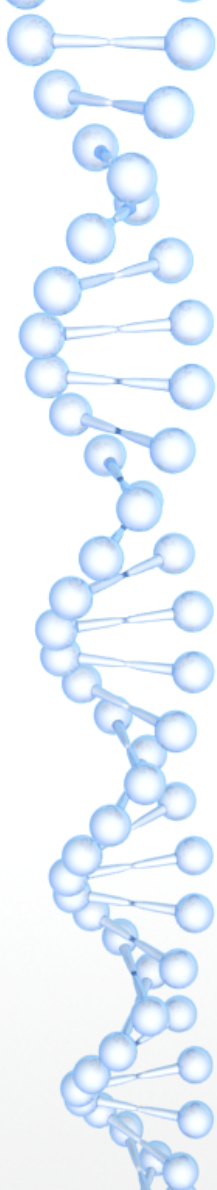


# Content of this talk

- Jupyter architecture overview
- Running a notebook remotely
- JupyterHub solution
- Alternative solution using remote desktops
- Pros and cons
- Future

# Jupyter architecture overview







# Running a notebook remotly

- Proxy requests
- Authenticate the users
- Central managment
- Share with other users
- Collaborative notebook editing
- Reconnect to a running notebook

# is it possible to make a jupyter notebook run even if the page is closed? #2446



laurentperrinet opened this issue 24 days ago · 1 comment



laurentperrinet commented 24 days ago



I am running multiple notebooks at the same time, often in different browsers, sometimes on different remote clients. till now, when I close the tab corresponding to a running notebook, it warns that the corresponding run will be stopped.

My question: How do I make a jupyter notebook resume it's run even if the page is closed ?

From what I understand, the client-server architecture could make that possible, but that there may be issues with multiple concurrent runs...

PS: I also asked the question on [SO](#) and will make sure to make this issue updated if I get solutions.



takluyver commented 24 days ago

Owner



Anything already running in the notebook will keep running, and the kernel it started for that will stay running - so it won't lose your variables. However, any output produced while the notebook isn't open in a browser tab is lost; there isn't an easy way to change this until we have the notebook server able to track the document state, which has been on the plan for ages.

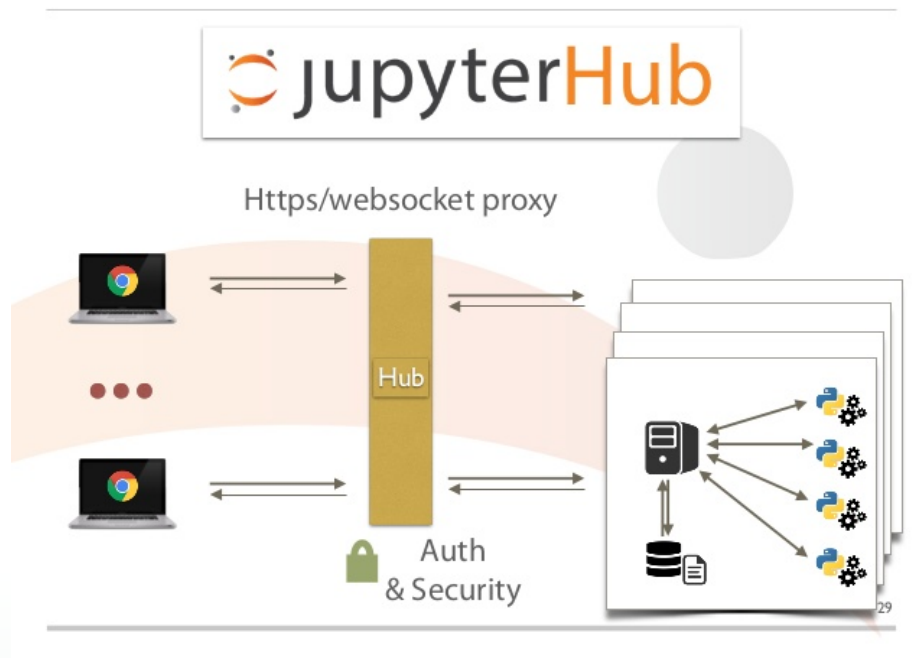


flying-sheep referenced this issue 3 days ago

**keep notebook running after the browser tab closed #1647**

**Closed**

# JupyterHub solution





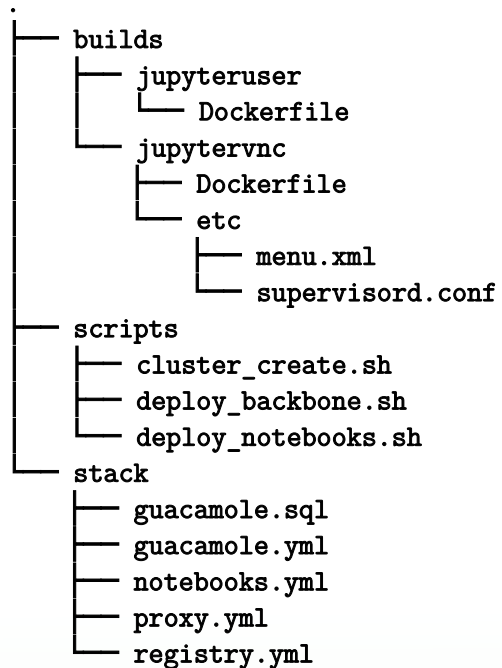


# Alternative solution

## Technologies used:

- Docker swarm cluster
- VNC server + Firefox + Jupyter
- Apache Guacamole (*clientless remote desktop gateway*)

# Alternative solution





# Docker cluster

```
# Create 3 machines
for i in 1 2 3; do
    docker-machine create -d virtualbox node-$i
done
eval $(docker-machine env node-1)

# Initialize docker manager
docker swarm init --advertise-addr $(docker-machine ip node-1)

# Join workers to the cluster
TOKEN=$(docker swarm join-token -q worker)
for i in 2 3; do
    eval $(docker-machine env node-$i)
    docker swarm join --token $TOKEN --advertise-addr $(docker-machine ip node-$i) \
        $(docker-machine ip node-1):2377
done
```



# Deploy backbone services

```
# Create proxy and notebooks networks
```

```
docker network create --driver overlay proxy
```

```
docker network create --driver overlay notebooks
```

```
# Start proxy, registry and guacamole
```

```
docker stack deploy -c ../stack/proxy.yml proxy
```

```
docker stack deploy -c ../stack/guacamole.yml guacamole
```

```
docker stack deploy -c ../stack/registry.yml registry
```

# Build and push images

- Common image (builds/jupyterenv/Dockerfile)

```
$ docker build -t localhost:5000/jupyterenv:5.0 .
```

```
$ docker push localhost:5000/jupyterenv:5.0
```

```
FROM debian:jessie-slim
RUN apt-get -y update \
    && apt-get install -y --no-install-recommends x11vnc xvfb supervisor openbox firefox-esr wget bzip2 ca-certificates \
    ...
# Configure environment
ENV CONDA_DIR /opt/conda
...
COPY etc/supervisord.conf /etc/supervisor/supervisord.conf
COPY etc/menu.xml /etc/xdg/openbox/menu.xml
...
RUN wget --quiet https://repo.continuum.io/miniconda/Miniconda3-4.3.11-Linux-x86_64.sh && \
    echo "b9fe70ce7b6fa8df05abfb56995959b897d0365299f5046063bc236843474fb8 *Miniconda3-4.3.11-Linux-x86_64.sh" | sha256sum -c - && \
    ...
USER $BG_USER
EXPOSE 5900
WORKDIR /home/$BG_USER
CMD ["/usr/bin/supervisord"]
```



# Build and push images

- User image (builds/jupyteruser/Dockerfile)

```
$ docker build -t localhost:5000/jupyteruser:5.0 .
```

```
$ docker push localhost:5000/jupyteruser:5.0
```

```
FROM localhost:5000/jupytervnc:5.0
```

```
ENV BG_NEW_USER username
```

```
ENV BG_NEW_UID 1391
```

```
# Change user UID
```

```
USER root
```

```
RUN usermod -u $BG_NEW_UID $BG_USER \
```

```
&& usermod -l $BG_NEW_USER $BG_USER \
```

```
&& chown -R -h $BG_NEW_USER /var/log/supervisor
```

```
USER $BG_NEW_USER
```



# Config and start a notebook

```
$ docker stack deploy -c ../stack/notebooks.yml nb
```

```
version: '3'

services:
  example:
    image: localhost:5000/jupyteruser:5.0
    networks:
      - notebooks
    volumes:
      - ${REPO_HOME}:/workspace
    working_dir: /workspace
    environment:
      - CONDA_ENVS_PATH=${REPO_HOME}/envs
```

# Config guacamole connection

## EDIT CONNECTION

Name:

Location:

Protocol:

## PARAMETERS

### Network

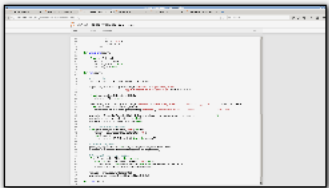
Hostname:

Port:

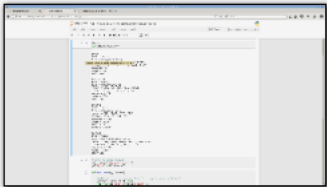


# Connect to the notebook

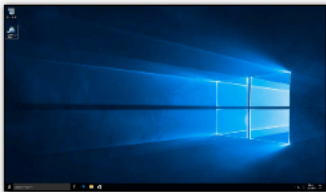
**RECENT CONNECTIONS** jdeu ▾



intogen



gendas

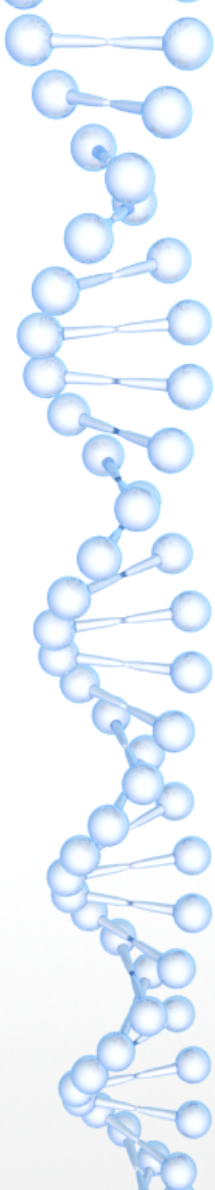


Windows 10

**ALL CONNECTIONS** Filter

- projects
  - gendas
  - intogen
- Windows 10

*Currently in use by 1 user.*



gendas x

bbglab.irbbarcelona.org/guacamole/#/client/MgBjAG15c3Fs

Rank scores - Mozilla Firefox

Rank scores x jdeu@2891e05eb7... x

localhost:8888/notebooks/src/examples/Rank scores.ipynb#

Jupyter Rank scores Last Checkpoint: Last Monday at 3:15 PM (autosaved)

File Edit View Insert Cell Kernel Help

Not Trusted Python [conda env:py35]

```
begin = BEGIN
end = END

In [4]: # Create a Gendas engine
gd = Gendas('data/gendas.conf')
gd['hg19'] = HG19Source()

In [5]: def mut_rank(gd, baserow):

    # Get scores of all the position in this gene group by tri>alt
    context = defaultdict(list)
    for r in gd['cadd'].merge(gd['hg19']):
        key = "{}>{}".format(r['hg19'][-1:1], r['cadd']['ALT'])
        context[key].append(r['cadd']['PHRED'])

    rows = []
    for m in gd['variants'].merge(gd['cadd'], on=['REF', 'ALT']).merge(gd['hg19']):

        key = "{}>{}".format(m['hg19'][-1:1], m['variants']['ALT'])
        ctx_scores = context[key]

        # Create the output row
        row = dict(baserow)
        for k, v in m['variants'].items():
            row[k] = v

        row['KEY'] = key
        row['SCORE'] = m['cadd']['PHRED']
        row['CONTEXT'] = ctx_scores

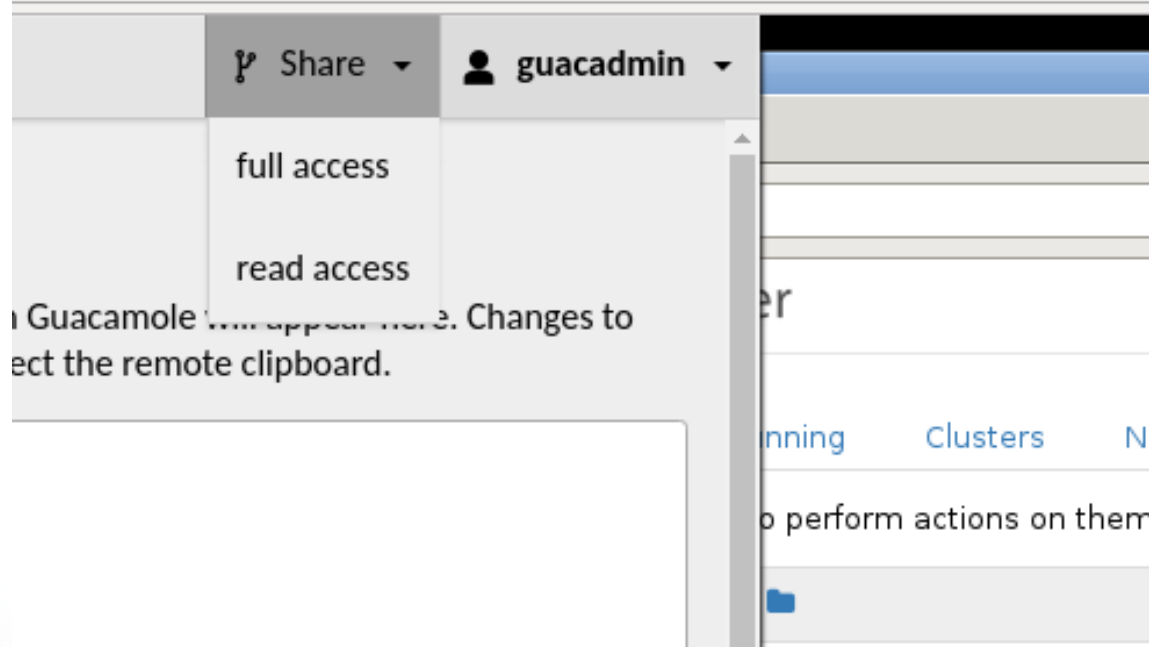
        rows.append(row)

    return rows

In [6]: %time
data = pd.DataFrame.from_dict(
    flatten(
```

# Share connection

58.99.100/guacamole/#/client/MQBjAG15c3Fs





# Pros and cons

- **Pros**
  - You can disconnect at any time
  - Concurrent notebook editing
  - Valid for any tool
- **Cons**
  - Indirect copy&paste
  - Scrolling less responsive
  - Different look&feel



## Future: to do list

- Add a GUI to notebook creation
- Integrate the solution in JupyterHub
- Or add a HTTP/Websockets proxy protocol in Apache Guacamole
- Create a Jupyter Lab image



# Future: collaborative editing

**Brian E. Granger, 6 Feb 2017**

Yes, we are doing this work in JupyterLab. We are building these features in a manner that will initially support Google Drive Real Time API, but could also have other real time backends plugged into it.

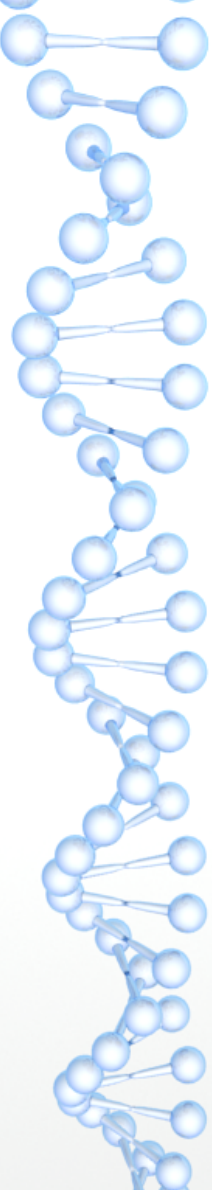
We have a full time post doc at UC Berkeley, Ian Rose, working on these things. The real time stuff isn't quite ready for public usage, but it is moving pretty fast at this point.


**<https://github.com/ian-r-rose/jupyterlab-google-drive>**

# Future: notebook hibernation

docker checkpoint create ...

*(already available with experimental enable)*





[Main page](#)  
[Recent changes](#)  
[Random page](#)  
[Help](#)

Tools

- [What links here](#)
- [Related changes](#)
- [Special pages](#)
- [Printable version](#)
- [Permanent link](#)
- [Page information](#)

News

- [Google+](#)
- [Twitter](#)


Page [Discussion](#)

## Docker

This article describes the status of CRIU integration with Docker, and how to use it.

**Contents** [\[hide\]](#)

- 1 Docker Experimental
  - 1.1 checkpoint
  - 1.2 restore
    - 1.2.1 Restoring into a **new** container
  - 1.3 Synopsis
- 2 Compatibility Notes
  - 2.1 TTY
  - 2.2 Seccomp
  - 2.3 OverlayFS
  - 2.4 Async IO
- 3 External checkpoint/restore



[github.com/jordeu/pytalks](https://github.com/jordeu/pytalks)

